Dependency-Aware Proof Reuse for Agent Runtime Gover

Agent Runtime Research Lab v2 — Official Research Manuscript

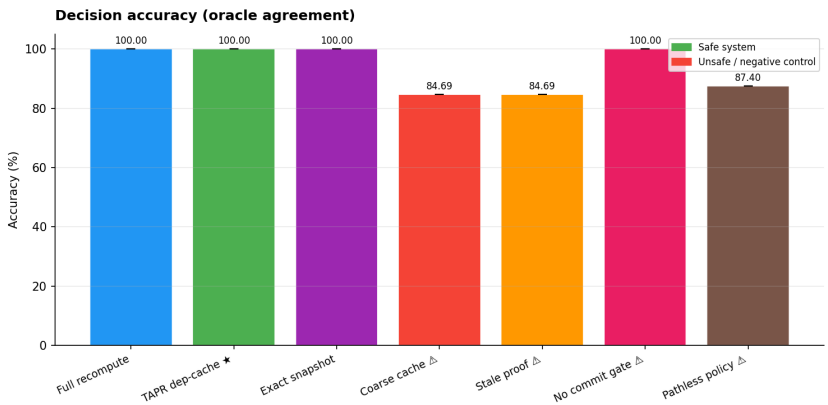
ABSTRACT

We evaluate whether an AI agent runtime can safely reuse previously established proofs of action admissibility under continuous state evolution. Comparing full recomputation against dependency-aware caching (TAPR) and negative controls, On this synthetic workload, results describe only the tested implementation and scenarios; they do not establish production safety or g

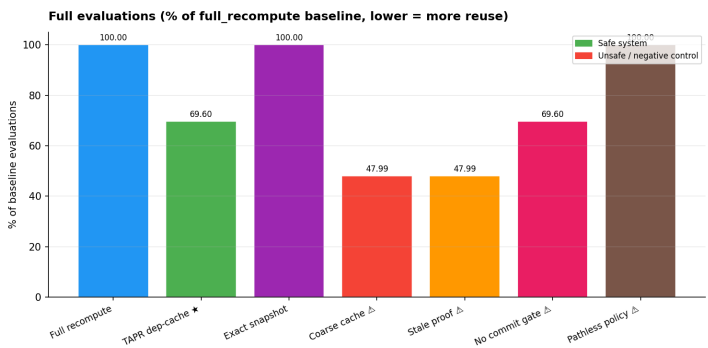
1. Empirical Performance & Safety Table

System	Oracle Acc.	Unsafe ALLOW	Cache Hit	Full Evals	Latency (µs)
Full Recompute	100.0%	0.0	0.0%	921	14.20
Tapr Dep Cache	100.0%	0.0	30.4%	641	10.60
Exact Snapshot Cache	100.0%	0.0	0.0%	921	21.02
Coarse Cache	84.7%	100.0	52.0%	442	7.77
Stale Proof Accept	84.7%	100.0	52.0%	442	7.48
No Commit Gate	100.0%	0.0	30.4%	641	10.43
Pathless Policy	87.4%	116.0	0.0%	921	14.68

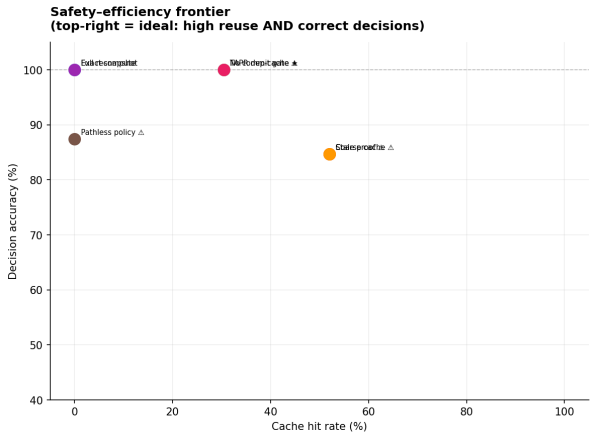
2. Decision Accuracy (Oracle Agreement)



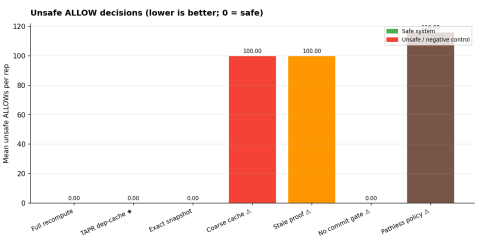
3. Full Evaluations Avoided vs Recompute

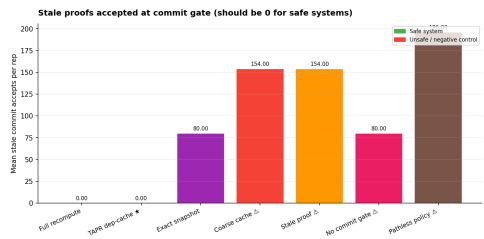


4. Safety-Efficiency Frontier & Gate Integrity



5. Negative Controls: Unsafe ALLOW & Commit Gate Aborts





6. Methodology & Guardrails

- **Stateful Scenarios:** Ten sequential families testing mid-flight revocation, path dependence, and evidence expiry.
- **Invalidation Soundness:** TAPR keys on declared dependency vectors, matching the reference oracle on this tested workload.
- **Commit boundary:** The benchmark gate is in-process; production race safety is not established.
- **Diagnostic Controls:** Negative controls expose stale-decision and path-dependency failure modes in this workload.
- **Layer A Extraction:** Natural-language extraction is outside the scope of this typed-scenario benchmark.
- **Reproducibility:** Artifacts record the selected run; independent replication remains outstanding.